

Machine Learning

Lecture 2: Linear models

Simon BERNARD

simon.bernard@univ-rouen.fr



Linear regression

First example: Hooker's experiment


- Joseph D. Hooker's experiment in 1849
- Measurement of atmospheric pressure p_i and boiling water temperature t_i in the Himalayas
- We know that $y_i = \ln(p_i)$ is (approx.) proportional to t_i , so we have :

$$y_i = \alpha t_i + \beta + u_i$$

where u_i represents the measurement error (or noise)

- Goal : Predict atmospheric pressure from the boiling temperature of water (which can then be used to deduce altitude, without using a barometer)

First example: Hooker's experiment



- Green points represent the measurements and the red line is the model $\hat{y}_i = wt_i + b$
- The parameter (w, b) are learned by minimizing the quadratic loss (least squares)
- This is called **simple linear regression** because there is only one input variable

Second example: Diabetes dataset ⁽¹⁾

Patient	AGE	SEX	BMI	BP	... Serum Measurements ...						Response
	x1	x2	x3	x4	x5	x6	x7	x8	x9	x10	y
1	59	2	32.1	101	157	93.2	38	4	4.9	87	151
2	48	1	21.6	87	183	103.2	70	3	3.9	69	75
3	72	2	30.5	93	156	93.6	41	4	4.7	85	141
4	24	1	25.3	84	198	131.4	40	5	4.9	89	206
5	50	1	23.0	101	192	125.4	52	4	4.3	80	135
6	23	1	22.6	89	139	64.8	61	2	4.2	68	97
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
441	36	1	30.0	95	201	125.2	42	5	5.1	85	220
442	36	1	19.6	71	250	133.2	97	3	4.6	92	57

Table 1. Diabetes study. 442 diabetes patients were measured on 10 baseline variables. A prediction model was desired for the response variable, a measure of disease progression one year after baseline.

- We assume that the relationship between the input features $x_{i,j}$ and the output y_i is linear:

$$y_i = w_1 x_{i,1} + w_2 x_{i,2} + \cdots + w_{10} x_{i,10} + b$$

- This is called **multiple linear regression** because there are multiple input variables

Second example: Diabetes dataset ⁽¹⁾

```
from sklearn import datasets
from sklearn.linear_model import LinearRegression
X, y = datasets.load_diabetes(scaled=False, return_X_y=True)
clf = LinearRegression().fit(X, y)
print(clf.coef_)
print(clf.intercept_)
```

```
> [-3.63612242e-02 -2.28596481e+01  5.60296209e+00  1.11680799e+00
    -1.08999633e+00  7.46450456e-01  3.72004715e-01  6.53383194e+00
     6.84831250e+01  2.80116989e-01]
-334.5671385187874
```

(1) Bradley Efron et al., "Least Angle Regression", Annals of Statistics, 2004

Matrix notations

- We have a set of n samples $(\mathbf{x}_i, \mathbf{y}_i), i = 1, \dots, n$ (the training set)
- In practice, the $\mathbf{x}_i$ are gathered in a matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$ and the $\mathbf{y}_i$ in a vector $\mathbf{y} \in \mathbb{R}^n$:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_i^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix} = \begin{bmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,j} & \dots & x_{1,d} \\ x_{2,1} & x_{2,2} & \dots & x_{2,j} & \dots & x_{2,d} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_{i,1} & x_{i,2} & \dots & x_{i,j} & \dots & x_{i,d} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_{n,1} & x_{n,2} & \dots & x_{n,j} & \dots & x_{n,d} \end{bmatrix}, \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_i \\ \vdots \\ y_n \end{bmatrix}$$

- Note : vectors are noted in bold, matrices in bold capitals, vectors are column vectors.

The least squares method

- We want to learn a linear model :

$$h(\mathbf{x}) = \sum_{j=1}^d w_j x_j + b = \mathbf{w}^\top \mathbf{x} + b$$

with $\mathbf{w}^\top = [w_1, w_2, \dots, w_d]$

- "linear" because it is linear in its parameters $\mathbf{w}$ and b (not necessarily in the x_j)
- The least squares method consists in finding the parameters $\mathbf{w}^*$ and b^* that minimize the quadratic loss, also called the residual sum of squares (RSS) :



$$(\mathbf{w}^*, b^*) = \arg \min_{\mathbf{w}, b} \text{RSS}(\mathbf{w}, b)$$

$$\text{RSS}(\mathbf{w}, b) = \sum_{i=1}^n (y_i - h(\mathbf{x}_i))^2 = \sum_{i=1}^n (y_i - \mathbf{w}^\top \mathbf{x}_i - b)^2$$

The least squares method

Geometric interpretation

The least squares method consists in finding the hyperplane that best fits the observations $(\mathbf{x}_i, y_i)$, such that the distances between the points and the hyperplane are minimized.



The least squares method

- For simplification, we usually add a column of 1 to $\mathbf{X}$ and incorporate b to the vector $\mathbf{w}$:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top & 1 \\ \mathbf{x}_2^\top & 1 \\ \vdots & \vdots \\ \mathbf{x}_n^\top & 1 \end{bmatrix} = \begin{bmatrix} x_{1,1} & \dots & x_{1,d} & 1 \\ x_{2,1} & \dots & x_{2,d} & 1 \\ \vdots & & \vdots & \vdots \\ x_{n,1} & \dots & x_{n,d} & 1 \end{bmatrix}, \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_d \\ b \end{bmatrix}$$

- The model becomes :

$$h(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

with $\mathbf{w}^\top = [w_1, w_2, \dots, w_d, b]$ and $\mathbf{x} = [x_1, x_2, \dots, x_d, 1]$

The least squares method

- The RSS can be written in matrix form :

$$\text{RSS}(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w})$$

- The minimum is reached when the gradient of the RSS is zero :

$$\nabla_{\mathbf{w}} \text{RSS}(\mathbf{w}) = -2\mathbf{X}^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

- This leads to the normal equations :

$$\mathbf{X}^\top \mathbf{X} \mathbf{w} = \mathbf{X}^\top \mathbf{y}$$

- If $\mathbf{X}^\top \mathbf{X}$ is invertible, the (**closed form**) solution is :

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

The least squares method

```
import numpy as np
from sklearn import datasets

X, y = datasets.load_diabetes(scaled=False, return_X_y=True)
Xa = np.concatenate((X, np.ones((X.shape[0], 1))), axis=1)
w = np.linalg.inv(Xa.T @ Xa) @ Xa.T @ y
print(w[:-1])
print(w[-1])
```

```
> [-3.63612242e-02 -2.28596481e+01  5.60296209e+00  1.11680799e+00
   -1.08999633e+00  7.46450456e-01  3.72004715e-01  6.53383194e+00
    6.84831250e+01  2.80116989e-01]
-334.5671385187874
```

Regularized regression

- The Least squares solution is :

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

- **Problem** : the matrix $\mathbf{X}^\top \mathbf{X}$ may not be invertible
- It may be the case if the features are linearly dependent (collinearity) or if there are too many features compared to the number of samples
- Solution : use a **regularization term to stabilize the inversion** :

$$\text{RSS}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^\top \mathbf{x}_i)^2 + \lambda \Omega(\mathbf{w})$$

- Optimization : solve the problem of $\mathbf{X}^\top \mathbf{X}$ not invertible
- Machine learning : aim to **penalize models that are too complex** to avoid overfitting and to favor the least complex solutions if several are possible

Ridge regression

- There exists several types of regularization, the most common is the **ridge regularization**
- The regularization term is the squared norm of the parameters $\mathbf{w}_j$ (called L_2 regularization) :

$$\text{RSSR}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \mathbf{w})^2 + \lambda \sum_{j=1}^d w_j^2$$

- The solution is given by :

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{S})^{-1} \mathbf{X}^\top \mathbf{y}$$

where

$$S_{i,j} = \begin{cases} 1 & \text{if } i = j \text{ and } i \leq d \\ 0 & \text{otherwise} \end{cases}$$


- λ is a hyperparameter that controls the "quantity" of regularization and is to be selected on a validation set or by cross-validation

Lasso regression

- The regularization term is the sum of the absolute values of the w_j (L_1 regularization) :

$$\text{RSSL}(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \mathbf{w})^2 + \lambda \sum_{j=1}^d |w_j|$$

- No closed-form solution, but can be solved with iterative algorithms
- Lasso regression has the property of producing sparse solutions, i.e. some w_j are exactly zero




- This allows to identify the most relevant features and to perform feature selection

Notebook `lec2-linear-regression.ipynb`

Linear classification

Linear classification


- Let's start with binary classification, i.e. $f : \mathbb{R}^d \rightarrow \mathcal{Y} = \{-1, 1\}$
- Linear model $h(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$
- How can we learn such a h to make it a good approximation of f ?
- First idea : use the least squares method and predict the class by taking the **sign** of $h(\cdot)$



Linear classification

Several drawbacks:

- No natural transposition to cases with > 2 classes
- Residuals (distance points-model) meaningless
- Result very sensitive to the training data and not always relevant



A posteriori probabilities

- Best alternative is to **predict probabilities of class membership**
- The linear model is learnt to predict $p(y = 1|\mathbf{x})$, called *a posteriori probability* of class 1 :

$$h(\mathbf{x}) = p(y = 1|\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

- For a binary classification with $y \in \{0, 1\}$, we have :


$$p(y = 0|\mathbf{x}) = 1 - p(y = 1|\mathbf{x}) = 1 - h(\mathbf{x})$$

- For all $\mathbf{x}$, the final prediction is 1 if $p(y = 1|\mathbf{x}) \geq p(y = 0|\mathbf{x})$, 0 otherwise

A posteriori probabilities

The main difficulty is that linear models are not suitable for modelling probabilities:

- can take any value in $\mathbb{R}$, while probabilities are in $[0, 1]$
- linear decision frontier, while probabilities are non-linear



Logistic regression

- Logistic regression is based on a **transformation of the output of the linear model**
- The model is learnt to predict the *logit* of the a posteriori probability of class 1 :

$$\text{logit}(p(y = 1|\mathbf{x})) = \mathbf{w}^\top \mathbf{x}$$

with

$$\text{logit}(p) = \ln \left(\frac{p}{1 - p} \right)$$

which gives :


$$\ln \left(\frac{p(y = 1|\mathbf{x})}{1 - p(y = 1|\mathbf{x})} \right) = \ln \left(\frac{p(y = 1|\mathbf{x})}{p(y = 0|\mathbf{x})} \right) = \mathbf{w}^\top \mathbf{x}$$

- A negative value means $p(y = 1|\mathbf{x}) < p(y = 0|\mathbf{x})$ and *vice versa*

Logistic regression

- We want to predict $p(y = 1|\mathbf{x})$ instead of $\text{logit}(p(y = 1|\mathbf{x}))$
- Use the reciprocal of *logit*, called **logistic function** or **sigmoid** :

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$




- The resulting logistic regression model is :

$$p(y = 1|\mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x})$$

Logistic regression

- The resulting logistic regression model is :

$$h(\mathbf{x}) = \sigma \left(\sum_{i=0}^d w_i x_i + b \right) = \sigma (\mathbf{w}^\top \mathbf{x})$$




Logistic regression

- The final prediction is given by $p(y = 1|\mathbf{x}) \underset{y=0}{\overset{y=1}{\geq}} p(y = 0|\mathbf{x})$
- The **decision frontier** is :

$$p(y = 1|\mathbf{x}) = p(y = 0|\mathbf{x}) = 0.5 = \sigma(\mathbf{w}^\top \mathbf{x})$$

that is, when $\mathbf{w}^\top \mathbf{x} = 0$



Loss minimization problem

- Parameter learning is based on the **maximum likelihood principle** ⁽¹⁾
- Results in the following loss function :

$$\begin{aligned} LL(\mathbf{w}) &= - \sum_{i=1}^n y_i \log(h(\mathbf{x}_i)) + (1 - y_i) \log(1 - h(\mathbf{x}_i)) \\ &= - \sum_{i=1}^n y_i \log(\sigma(\mathbf{w}^\top \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^\top \mathbf{x}_i)) \end{aligned}$$

with $y_i \in \{0, 1\}$ the class of the i -th sample

- This loss function is known as the **binary cross-entropy loss**

(1) This principle is not explained here as it is beyond the scope of this course.

Loss minimization problem

- The underlying minimization problem is :


$$\mathbf{w}^* = \arg \min_{\mathbf{w}} LL(\mathbf{w})$$

- Same approach as usual : calculate the gradient and set it to 0

$$\nabla LL(\mathbf{w}) = \mathbf{X}^\top (\sigma(\mathbf{X}\mathbf{w}) - \mathbf{y}) = \mathbf{X}^\top (h(\mathbf{X}) - \mathbf{y}) = 0$$

- We'll spare you the math — it does not result in a closed form solution.
- Solution : iterative numerical method called **gradient descent**.

Gradient descent



1. Initialise $\mathbf{w}_0$ with random values
2. Calculate the gradient $\nabla LL(\mathbf{w}_0)$
3. Update the parameters using the gradient :

$$\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k - \eta \nabla LL(\mathbf{w}_k)$$

with η the **learning rate** (hyperparameter)

4. Repeat step 3 until convergence

Gradient descent

- Learning rate η crucial to control the convergence speed
 - Too small $\Rightarrow$ slow convergence; too large $\Rightarrow$ may diverge
 - Strategies to adapt η during learning (e.g. backtracking line search)
- Reaching convergence is not trivial
 - Depends on η , parameter initialization, loss function, etc.
 - Stopping criterion : ϵ usually set to an arbitrarily small value (e.g. 10^{-6})
 - Alternative convergence criteria (e.g. based on the norm of the gradient)
- Well-known gradient descent based optimizers : Newton GD, Adam, RMSProp, etc. ⁽²⁾

(2) Not explained in this course, but you might encounter them in practice

Notebook `lec2-logistic-regression.ipynb`

Stochastic Gradient Descent

- Empirical risk minimization setup :


$$LL(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y_i, h_{\mathbf{w}}(\mathbf{x}_i))$$
$$\nabla LL(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\mathbf{w}} \mathcal{L}(y_i, h_{\mathbf{w}}(\mathbf{x}_i))$$

- Complexity of an update grows linearly with n : **bad idea when dealing with large datasets**
- **Stochastic Gradient Descent (SGD)** : update using only one training instance at a time


$$\mathbf{w}_{k+1} = \mathbf{w}_k - \eta \nabla LL(\mathbf{w}_k)$$

$$LL(\mathbf{w}_k) = \mathcal{L}(y_{rand}, h_{\mathbf{w}_k}(\mathbf{x}_{rand}))$$


Stochastic Gradient Descent



Batch Gradient Descent



Mini-batch Gradient Descent



Stochastic Gradient Descent

Stochastic Gradient Descent

- With very large dataset : trade-off between Batch GD and Stochastic GD
- SGD results in more updates but with raw approximation of the gradient
- BGD results in fewer updates but with better approximation of the gradient
- What is the best?
- Usually more updates is better but with really big datasets, convergence is hard to reach
- In practice, computational costs often dictate the size of the batch

Softmax regression

- Reminder for binary classification : $y \in \{0, 1\}$
- Predict *a posteriori* probabilities :

$$p(y = 0|\mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(\mathbf{w}^\top \mathbf{x})}$$

$$p(y = 1|\mathbf{x}) = 1 - p(y = 0|\mathbf{x}) = \frac{\exp(\mathbf{w}^\top \mathbf{x})}{1 + \exp(\mathbf{w}^\top \mathbf{x})}$$

- Decision rule : $\hat{y} = 0$ if $p(y = 0|\mathbf{x}) \geq p(y = 1|\mathbf{x})$, $\hat{y} = 1$ otherwise

Softmax regression

- Multiclass classification : $y \in 1, \dots, K, K > 2$
- Predict *a posteriori* probabilities :


$$p(y = k|\mathbf{x}) = \frac{\exp(\mathbf{w}_k^\top \mathbf{x})}{\sum_{l=1}^K \exp(\mathbf{w}_l^\top \mathbf{x})}, \quad \forall k = 1, \dots, K$$

where $\mathbf{w}_k$ is the vector of parameters for class k

- This is the **softmax function**, a generalization of the sigmoid function
- It allows to have a probability distribution over the K classes (sums to 1)
- Decision rule : $\hat{y} = \arg \max_k p(y = k|\mathbf{x})$

Softmax regression

Example : 2D version of Iris dataset



Takeaways

Linear models are simple and interpretable, but may not capture complex patterns

Linear models are foundational for deep learning and many other machine learning techniques

Regularization is a key concept, and is often essential to prevent overfitting and improve generalization

Gradient descent and its variants are the most used technique for optimizing models, especially in deep learning

The softmax function is a key component for the design of deep learning methods for classification tasks, allowing to model probabilities over multiple classes